


Validation in a separate higher-fidelity simulator (tabula_drone: spatial, HP-depletion, episodic): SwarmCF beats the simulator's SGD-MF and Independent-UCB, approaching the oracle

$\text{skill} = (\text{policy} - \text{random}) / (\text{oracle} - \text{random})$

Converged skill



Learning curves

